




Building Better Agents with Less Context


RECURSIVE SKILL BUILDING & PRODUCTIVITY

A mechanical breakdown of how AI processes information, and why Recursive Skill Building is the only reliable way to scale for productivity.

THE BOTTLENECK IS THE HARNESS, NOT THE INTELLIGENCE.


Today's frontier models are exceptionally good. If your agent is failing, the problem isn't a lack of AI intelligence. The problem is your workflow expectations. We don't lack intelligence—we lack context management.

WHAT PEOPLE BUILD (SCALING FOR COOL)



What people build (Scaling for Cool)

WHAT ACTUALLY WORKS (SCALING FOR PRODUCTIVITY)



What actually works (Scaling for Productivity)

MODELS PREDICT TOKENS. THEY DO NOT THINK.

AI does not possess reasoning or understanding. When you ask for a report, it maps your English to the closest mathematical resemblance on a vector graph. Without a codified step-by-step guide, it is merely guessing.


VECTOR GRAPH



THE ANATOMY OF THE CONTEXT WINDOW


Every piece of information an AI needs to execute an action must be assembled here. Space is strictly finite, and every token costs processing power and accuracy.

**10% IS ALREADY FULL
BEFORE YOU START.**




THE CONTEXT TRAP: WHY LESS IS MORE

As the context window approaches 100%, the model gets objectively dumb and forgets instructions. Forcing heavy global instructions on every single conversational turn degrades performance and wastes money.



STOP FORCING GLOBAL CONTEXT WITH THE 95% RULE.

95% of people do not need a massive agent.md file. Reserve global files strictly for proprietary, company-wide data that must be referenced on literally every turn. For everything else, rely on modular skills.




Diagnostic Matrix: Global Files vs. Modular Skills

	<u>agent.md</u>	<u>skill.md</u>
Mechanism	Forces load on every single turn	Loads on demand (Progressive disclosure)
Token Cost	High (Example: 944 tokens per run for 116 lines)	Low (Example: 53 tokens until activated)
Best Use Case	Proprietary, universal company data	Specific workflows (e.g., code structure, reports)
Risk Profile	Rapidly clutters the context window 	Highly performant, preserves context 

The invisible leverage of Progressive Disclosure.

Skills keep your context window pristine. The agent holds only the title and description in its memory until the exact moment of execution.



Three critical errors in agent building.



Don't build zero-context skills.

Writing a skill before establishing what a successful manual run looks like guarantees failure.



Don't download marketplace skills.

Generic skills lack your specific workflow nuances and open you up to external attack vectors.




Don't expect immediate perfection.

A skill will not work on the first try. Expect 3 to 6 hiccups as the baseline cost of setup.


The Recursive Skill Building Framework

Stop treating agents like black magic boxes. Treat them like capable new hires who need a baseline rundown of your process. Establish success manually before you automate.




Phase 1: Identify & Walkthrough

Pinpoint a repeatable workflow. Do not write the code. Run the task live with the agent. Let it fail. Correct it manually to establish exactly what a 100% successful run looks like strictly within the chat context.




Phase 2: AI Generates the Skill

Once you have a perfect manual run, command the AI to review its own successful context and write the skill for you. The AI codifies the exact steps and nuances required. You never have to handwrite a complex skill again.




Phase 3: The Error Feedback Loop

When the skill fails on an edge case, do not get frustrated. Pass the exact failure back to the agent. Ask it to fix the issue, and command it to update the skill file. Five loops will yield flawless execution.



Scale for productivity, not for what looks cool.

Do not spin up 15 sub-agents on day one. Build a deeply productive stack by starting with one agent, recursively building your personal skills, and only adding specialized sub-agents once the foundation is rock solid.



Your ultimate moat is your specific workflow.

General knowledge is now a commodity. Frontier models will only get smarter. Your protection is your specific taste, your specific strategy, and your ability to codify them into perfectly tuned AI skills. Build your own.

